

Phase 3 :- Execution

Problem Statement :-

Air Quality Analysis and Prediction in Tamilnadu

The question involves analyzing and predicting air quality in the state of Tamil Nadu. Specifically, it focuses on understanding the relationship between air quality parameters, such as sulfur dioxide (SO₂) and nitrogen dioxide (NO₂), and particulate matter (RSPM/PM₁₀), which can impact air quality and human health.

Data set Details:-

You obtained the dataset from Kaggle. The dataset can be found at this link: [Air Quality Data Set](#). It contains information about air quality measurements in Tamil Nadu in 2014.

Data set Info and Description:-

- **SO₂ (Sulfur Dioxide):** SO₂ is a gaseous air pollutant produced by the burning of fossil fuels, particularly in industrial processes. It is a key contributor to air pollution and can have adverse health effects.
- **NO₂ (Nitrogen Dioxide):** NO₂ is another gaseous air pollutant, often associated with vehicle emissions and industrial processes. Like SO₂, it can also impact air quality and health.
- **RSPM/PM₁₀ (Particulate Matter):** Particulate matter refers to tiny solid particles or liquid droplets in the air. PM₁₀ specifically refers to particles with a diameter of 10 micrometers or smaller. These particles can originate from various sources and affect air quality and respiratory health.

Libraries used :-

You used several libraries in your Jupyter notebook, including **pandas** for data manipulation, **numpy** for numerical operations,

matplotlib.pyplot for data visualization, and **sklearn** (scikit-learn) for machine learning tools.


Data Preprocessing:-

- You renamed the column 'RSPM/PM10' to 'RSPMorPM10' using **df.rename(columns={'RSPM/PM10': 'RSPMorPM10'})** for easier reference.
- You checked for missing values in the DataFrame using **df.isnull()** and calculated the total number of missing values using **.sum().sum()**. There were 2907 missing values in the dataset.
- You filled the missing values with zeros using **df.fillna(0)** to ensure all data points are numeric.

Data Visualization:-

You created a new DataFrame `cdf` containing a subset of columns: 'SO2', 'NO2', and 'RSPMorPM10'.

You visualized the relationship between 'SO2' (Sulfur Dioxide) and 'RSPMorPM10' (Particulate Matter) using a scatter plot. This plot helps you visualize the data points and the potential relationship between the two variables.



Training and Testing:-

In your Jupyter notebook, you followed these steps:

- You split the dataset into a training set and a testing set using a random mask.
- You created a linear regression model using scikit-learn's **`linear_model.LinearRegression()`**.
- You trained the model on the training data using 'SO2' and 'NO2' as features and 'RSPMorPM10' as the target variable.
- You made predictions on the test data using the trained model.
- You evaluated the model's performance by calculating the Mean Squared Error (MSE) and the Variance score (R^2) to check how well the model predicts 'RSPMorPM10' based on 'SO2' and 'NO2'.

Analysis:-

The rest of your analysis involved data preprocessing, including renaming columns and handling missing values, data visualization to understand the relationship between 'SO2' and 'RSPMorPM10', and building and evaluating a linear regression model for predicting 'RSPMorPM10'.

Model Evaluation:-


- You made predictions on the test data using the trained model with **`y_hat = regr.predict(test[['SO2', 'NO2']])`**. These predictions represent the model's estimates of 'RSPMorPM10' based on 'SO2' and 'NO2'.
- You calculated the Mean Squared Error (MSE) using **`np.mean((y_hat - y) ** 2)`**. The MSE measures the average squared difference between the predicted and actual values. A lower MSE indicates a better-performing model.

- You also calculated the Variance score (R^2) using **regr.score(x, y)**. The R^2 score measures the proportion of the variance in the dependent variable ('RSPMorPM10') that is predictable from the independent variables ('SO2' and 'NO2'). A higher R^2 score (close to 1) indicates a better fit of the model to the data. In this case, the R^2 score is approximately 0.17, suggesting that the model explains only a small portion of the variance in 'RSPMorPM10'.

Metrics Used :-

- For accuracy check, you used the Mean Squared Error (MSE) and the Variance score (R^2):
 - **Mean Squared Error (MSE):** This metric measures the average squared difference between the predicted and actual values. A lower MSE indicates a better-performing model. In your analysis, the MSE was approximately 720.36, suggesting the model's predictions were not very close to the actual values on average.
 - **Variance score (R^2):** This score represents the proportion of the variance in the dependent variable ('RSPMorPM10') that is predictable from the independent variables ('SO2' and 'NO2'). A higher R^2 score (close to 1) indicates a better fit of the model to the data. In your analysis, the R^2 score was approximately 0.17, indicating that the model explains only a small portion of the variance in 'RSPMorPM10'.

SIMPLE LINEAR REGRESSION :



```
In [24]: from sklearn import linear_model
regr = linear_model.LinearRegression()
train = train.dropna()
x = np.asanyarray(train[['SO2', 'NO2']])
y = np.asanyarray(train[['RSPMorPM10']])
regr.fit(x, y)
# The coefficients
print ('Coefficients: ', regr.coef_)

Coefficients:  [[2.80336749  0.18971444]]
```


```
In [25]: test = test.dropna()
y_hat= regr.predict(test[['SO2','NO2']])
x = np.asanyarray(test[['SO2','NO2']])
y = np.asanyarray(test[['RSPMorPM10']])
print("Mean Squared Error (MSE) : %.2f"
% np.mean((y_hat - y) ** 2))
# Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % regr.score(x, y))

Mean Squared Error (MSE) : 794.16
Variance score: 0.19
```

RANDOM FOREST ALGORITHM :

Mean Squared Error: 775.49

R-squared (R2) Score: 0.25



```
# Handle missing values by filling with the mean
data = data.fillna(data.mean())

# Select the relevant features (SO2 and NO2) and the target variable (RSPM/PM10)
features = data[["SO2", "NO2"]]
target = data["RSPM/PM10"]

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(features, target, test_size=0.2, random_state=42)

# Create and train the Random Forest model
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = rf_model.predict(X_test)

# Calculate model performance metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)


print(f"Mean Squared Error: {mse:.2f}")
print(f"R-squared (R2) Score: {r2:.2f}")

# Visualize the feature importances
feature_importances = rf_model.feature_importances_
plt.bar(features.columns, feature_importances)
plt.xlabel("Features")
plt.ylabel("Importance")
plt.title("Feature Importances")
plt.show()
```

REGRESSION WITH GRADIENT BOOSTING ALGORITHM :

Mean Squared Error: 684.02

R-squared (R2) Score: 0.34



```
# Handle missing values by filling with the mean
data = data.fillna(data.mean())

# Select the relevant features (SO2 and NO2) and the target variable (RSPM/PM10)
features = data[["SO2", "NO2"]]
target = data["RSPM/PM10"]

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(features, target, test_size=0.2, random_state=42)

# Create and train the Gradient Boosting model
gb_model = GradientBoostingRegressor(n_estimators=100, random_state=42)
gb_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = gb_model.predict(X_test)

# Calculate model performance metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse:.2f}")
print(f"R-squared (R2) Score: {r2:.2f}")


# Visualize the feature importances
feature_importances = gb_model.feature_importances_
plt.bar(features.columns, feature_importances)
plt.xlabel("Features")
plt.ylabel("Importance")
plt.title("Feature Importances")
plt.show()
```

GAUSSIAN PROCESS REGRESSION (GPR) :

RMSE: 33.17

R-squared: -0.05

Gaussian Process Regression - Actual vs. Predicted



```
# Standardize the features
scaler = StandardScaler()
X = scaler.fit_transform(X)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Define the Gaussian Process Kernel
kernel = C(1.0, (1e-3, 1e3)) * RBF(1.0, (1e-2, 1e2))

# Create and train the Gaussian Process Regressor
gpr = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=10, random_state=42)
gpr.fit(X_train, y_train)

# Make predictions
y_pred, sigma = gpr.predict(X_test, return_std=True)

# Evaluate the model (you can use various metrics, e.g., RMSE, R-squared)
from sklearn.metrics import mean_squared_error, r2_score
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
r2 = r2_score(y_test, y_pred)
print(f'RMSE: {rmse:.2f}')
print(f'R-squared: {r2:.2f}')

# Visualize the results
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, color='blue', alpha=0.6)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], linestyle='--', color='red', linewidth=2)
plt.xlabel('Actual RSPM/PM10')
plt.ylabel('Predicted RSPM/PM10')
plt.title('Gaussian Process Regression - Actual vs. Predicted')
plt.show()
```


We've loaded the dataset, performed data preprocessing, and explored different machine learning algorithms to predict air quality (RSPM/PM10) based on the levels of SO2 and NO2. We've tried linear regression, random forest, gradient boosting, and Gaussian process regression. Here's a summary of the results for each model:

1. ***Linear Regression:***

- Mean Squared Error (MSE): 794.16
- R-squared (R2) Score: 0.19
- Interpretation: The linear regression model didn't perform well, as indicated by the low R-squared score.

2. ***Random Forest Regression:***

- Mean Squared Error (MSE): 775.49
- R-squared (R2) Score: 0.25
- Interpretation: The random forest model performed slightly better than linear regression but still has room for improvement.

3. ***Gradient Boosting Regression:***

- Mean Squared Error (MSE): 684.02
- R-squared (R2) Score: 0.34
- Interpretation: The gradient boosting model improved the results compared to the previous models.

4. ***Gaussian Process Regression (GPR):***

- Gaussian Process Regression is another technique you've applied, but there's no specific performance metric provided. GPR can capture complex relationships but requires careful kernel selection and tuning.

Overall, the best-performing model among the ones we've tried is the Gradient Boosting Regression, with the highest R-squared score and the lowest Mean Squared Error.

Importing Libraries

```
import matplotlib.pyplot as plt
import pandas as pd
import pylab as pl
import numpy as np
import sklearn
%matplotlib inline
```

Loading the Data Set

```
df = pd.read_csv("cpcb_dly_aq_tamil_nadu-2014.csv")
df = df.rename(columns={'RSPM/PM10': 'RSPMorPM10'})
df
```

	Stn Code	Sampling Date	State	City/Town/Village/Area	\
0	38	01-02-2014	Tamil Nadu	Chennai	
1	38	01-07-2014	Tamil Nadu	Chennai	
2	38	21-01-2014	Tamil Nadu	Chennai	
3	38	23-01-2014	Tamil Nadu	Chennai	
4	38	28-01-2014	Tamil Nadu	Chennai	
...
2874	773	12-03-2014	Tamil Nadu	Trichy	
2875	773	12-10-2014	Tamil Nadu	Trichy	
2876	773	17-12-2014	Tamil Nadu	Trichy	
2877	773	24-12-2014	Tamil Nadu	Trichy	
2878	773	31-12-2014	Tamil Nadu	Trichy	

	Location of Monitoring Station	\
0	Kathivakkam, Municipal Kalyana Mandapam, Chennai	
1	Kathivakkam, Municipal Kalyana Mandapam, Chennai	
2	Kathivakkam, Municipal Kalyana Mandapam, Chennai	
3	Kathivakkam, Municipal Kalyana Mandapam, Chennai	
4	Kathivakkam, Municipal Kalyana Mandapam, Chennai	
...
2874	Central Bus Stand, Trichy	
2875	Central Bus Stand, Trichy	
2876	Central Bus Stand, Trichy	
2877	Central Bus Stand, Trichy	
2878	Central Bus Stand, Trichy	

	Agency	\
0	Tamilnadu State Pollution Control Board	
1	Tamilnadu State Pollution Control Board	
2	Tamilnadu State Pollution Control Board	
3	Tamilnadu State Pollution Control Board	
4	Tamilnadu State Pollution Control Board	
...
2874	Tamilnadu State Pollution Control Board	

```

2875 Tamilnadu State Pollution Control Board
2876 Tamilnadu State Pollution Control Board
2877 Tamilnadu State Pollution Control Board
2878 Tamilnadu State Pollution Control Board

Type of Location    S02    N02    RSPMorPM10    PM
2.5
0      Industrial Area    11.0    17.0          55.0
NaN
1      Industrial Area    13.0    17.0          45.0
NaN
2      Industrial Area    12.0    18.0          50.0
NaN
3      Industrial Area    15.0    16.0          46.0
NaN
4      Industrial Area    13.0    14.0          42.0
NaN
...      ...      ...      ...      ...
..
2874 Residential, Rural and other Areas    15.0    18.0          102.0
NaN
2875 Residential, Rural and other Areas    12.0    14.0          91.0
NaN
2876 Residential, Rural and other Areas    19.0    22.0          100.0
NaN
2877 Residential, Rural and other Areas    15.0    17.0          95.0
NaN
2878 Residential, Rural and other Areas    14.0    16.0          94.0
NaN

[2879 rows x 11 columns]

```

Exploring the Data Set

```
print(df.head())
```

```

   Stn Code Sampling Date      State City/Town/Village/Area \
0        38    01-02-2014  Tamil Nadu                Chennai
1        38    01-07-2014  Tamil Nadu                Chennai
2        38    21-01-2014  Tamil Nadu                Chennai
3        38    23-01-2014  Tamil Nadu                Chennai
4        38    28-01-2014  Tamil Nadu                Chennai

   Location of Monitoring Station \
0  Kathivakkam, Municipal Kalyana Mandapam, Chennai
1  Kathivakkam, Municipal Kalyana Mandapam, Chennai
2  Kathivakkam, Municipal Kalyana Mandapam, Chennai
3  Kathivakkam, Municipal Kalyana Mandapam, Chennai
4  Kathivakkam, Municipal Kalyana Mandapam, Chennai

```

	Agency	Type of Location	S02	N02
0	Tamilnadu State Pollution Control Board	Industrial Area	11.0	17.0
1	Tamilnadu State Pollution Control Board	Industrial Area	13.0	17.0
2	Tamilnadu State Pollution Control Board	Industrial Area	12.0	18.0
3	Tamilnadu State Pollution Control Board	Industrial Area	15.0	16.0
4	Tamilnadu State Pollution Control Board	Industrial Area	13.0	14.0

	RSPMorPM10	PM 2.5
0	55.0	NaN
1	45.0	NaN
2	50.0	NaN
3	46.0	NaN
4	42.0	NaN

```
print(df.tail())
```

	Stn Code	Sampling Date	State	City/Town/Village/Area
2874	773	12-03-2014	Tamil Nadu	Trichy
2875	773	12-10-2014	Tamil Nadu	Trichy
2876	773	17-12-2014	Tamil Nadu	Trichy
2877	773	24-12-2014	Tamil Nadu	Trichy
2878	773	31-12-2014	Tamil Nadu	Trichy

	Location of Monitoring Station	Agency
2874	Central Bus Stand, Trichy	Tamilnadu State Pollution Control Board
2875	Central Bus Stand, Trichy	Tamilnadu State Pollution Control Board
2876	Central Bus Stand, Trichy	Tamilnadu State Pollution Control Board
2877	Central Bus Stand, Trichy	Tamilnadu State Pollution Control Board
2878	Central Bus Stand, Trichy	Tamilnadu State Pollution Control Board

	Type of Location	S02	N02	RSPMorPM10	PM 2.5
2874	Residential, Rural and other Areas	15.0	18.0	102.0	NaN
2875	Residential, Rural and other Areas	12.0	14.0	91.0	NaN
2876	Residential, Rural and other Areas	19.0	22.0	100.0	NaN

2877	Residential, Rural and other Areas	15.0	17.0	95.0
NaN				
2878	Residential, Rural and other Areas	14.0	16.0	94.0
NaN				

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 2879 entries, 0 to 2878
```

```
Data columns (total 11 columns):
```

#	Column	Non-Null Count	Dtype
0	Stn Code	2879 non-null	int64
1	Sampling Date	2879 non-null	object
2	State	2879 non-null	object
3	City/Town/Village/Area	2879 non-null	object
4	Location of Monitoring Station	2879 non-null	object
5	Agency	2879 non-null	object
6	Type of Location	2879 non-null	object
7	S02	2868 non-null	float64
8	N02	2866 non-null	float64
9	RSPMorPM10	2875 non-null	float64
10	PM 2.5	0 non-null	float64

```
dtypes: float64(4), int64(1), object(6)
```

```
memory usage: 247.5+ KB
```

```
None
```

```
print(df.describe())
```

	Stn Code	S02	N02	RSPMorPM10	PM 2.5
count	2879.000000	2868.000000	2866.000000	2875.000000	0.0
mean	475.750261	11.503138	22.136776	62.494261	NaN
std	277.675577	5.051702	7.128694	31.368745	NaN
min	38.000000	2.000000	5.000000	12.000000	NaN
25%	238.000000	8.000000	17.000000	41.000000	NaN
50%	366.000000	12.000000	22.000000	55.000000	NaN
75%	764.000000	15.000000	25.000000	78.000000	NaN
max	773.000000	49.000000	71.000000	269.000000	NaN

Identifying null Values

```
print(df.isnull())
```

	Stn Code	Sampling Date	State	City/Town/Village/Area	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	
...	

2874	False	False	False	False
2875	False	False	False	False
2876	False	False	False	False
2877	False	False	False	False
2878	False	False	False	False

	Location of Monitoring Station	Agency	Type of Location	S02
N02 \				
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False
...
2874	False	False	False	False
2875	False	False	False	False
2876	False	False	False	False
2877	False	False	False	False
2878	False	False	False	False

	RSPMorPM10	PM 2.5
0	False	True
1	False	True
2	False	True
3	False	True
4	False	True
...
2874	False	True
2875	False	True
2876	False	True
2877	False	True
2878	False	True

[2879 rows x 11 columns]

```
c = df.isnull().sum()
print(c)
```

```

Stn Code                0
Sampling Date           0
State                  0
City/Town/Village/Area 0
Location of Monitoring Station 0
Agency                0
Type of Location       0
SO2                    11
NO2                    13
RSPMorPM10             4
PM 2.5                 2879
dtype: int64

print('Total Sum of null values in the Data set = ',c.sum())

Total Sum of null values in the Data set = 2907

print(df['RSPMorPM10'].value_counts()) #frequency of values

47.0      64
41.0      62
43.0      59
51.0      58
40.0      58
..
163.0     1
138.0     1
211.0     1
202.0     1
238.0     1
Name: RSPMorPM10, Length: 169, dtype: int64

```

Data Preprocessing - Replacing the null values

```

df.drop_duplicates()

```

	Stn Code	Sampling Date	State	City/Town/Village/Area	\
0	38	01-02-2014	Tamil Nadu	Chennai	
1	38	01-07-2014	Tamil Nadu	Chennai	
2	38	21-01-2014	Tamil Nadu	Chennai	
3	38	23-01-2014	Tamil Nadu	Chennai	
4	38	28-01-2014	Tamil Nadu	Chennai	
...	
2874	773	12-03-2014	Tamil Nadu	Trichy	
2875	773	12-10-2014	Tamil Nadu	Trichy	
2876	773	17-12-2014	Tamil Nadu	Trichy	
2877	773	24-12-2014	Tamil Nadu	Trichy	
2878	773	31-12-2014	Tamil Nadu	Trichy	
					Location of Monitoring Station \

0	Kathivakkam, Municipal	Kalyana Mandapam, Chennai
1	Kathivakkam, Municipal	Kalyana Mandapam, Chennai
2	Kathivakkam, Municipal	Kalyana Mandapam, Chennai
3	Kathivakkam, Municipal	Kalyana Mandapam, Chennai
4	Kathivakkam, Municipal	Kalyana Mandapam, Chennai
...		...
2874		Central Bus Stand, Trichy
2875		Central Bus Stand, Trichy
2876		Central Bus Stand, Trichy
2877		Central Bus Stand, Trichy
2878		Central Bus Stand, Trichy

		Agency \
0	Tamilnadu State Pollution Control Board	
1	Tamilnadu State Pollution Control Board	
2	Tamilnadu State Pollution Control Board	
3	Tamilnadu State Pollution Control Board	
4	Tamilnadu State Pollution Control Board	
...		...
2874	Tamilnadu State Pollution Control Board	
2875	Tamilnadu State Pollution Control Board	
2876	Tamilnadu State Pollution Control Board	
2877	Tamilnadu State Pollution Control Board	
2878	Tamilnadu State Pollution Control Board	

	Type of Location	S02	N02	RSPMorPM10	PM
2.5					
0	Industrial Area	11.0	17.0	55.0	
NaN					
1	Industrial Area	13.0	17.0	45.0	
NaN					
2	Industrial Area	12.0	18.0	50.0	
NaN					
3	Industrial Area	15.0	16.0	46.0	
NaN					
4	Industrial Area	13.0	14.0	42.0	
NaN					
...
..					
2874	Residential, Rural and other Areas	15.0	18.0	102.0	
NaN					
2875	Residential, Rural and other Areas	12.0	14.0	91.0	
NaN					
2876	Residential, Rural and other Areas	19.0	22.0	100.0	
NaN					
2877	Residential, Rural and other Areas	15.0	17.0	95.0	
NaN					
2878	Residential, Rural and other Areas	14.0	16.0	94.0	
NaN					

[2879 rows x 11 columns]

df.fillna(0)

	Stn	Code	Sampling Date	State	City/Town/Village/Area	\
0		38	01-02-2014	Tamil Nadu		Chennai
1		38	01-07-2014	Tamil Nadu		Chennai
2		38	21-01-2014	Tamil Nadu		Chennai
3		38	23-01-2014	Tamil Nadu		Chennai
4		38	28-01-2014	Tamil Nadu		Chennai
...
2874		773	12-03-2014	Tamil Nadu		Trichy
2875		773	12-10-2014	Tamil Nadu		Trichy
2876		773	17-12-2014	Tamil Nadu		Trichy
2877		773	24-12-2014	Tamil Nadu		Trichy
2878		773	31-12-2014	Tamil Nadu		Trichy

	Location of Monitoring Station					\
0	Kathivakkam, Municipal	Kalyana	Mandapam,	Chennai		
1	Kathivakkam, Municipal	Kalyana	Mandapam,	Chennai		
2	Kathivakkam, Municipal	Kalyana	Mandapam,	Chennai		
3	Kathivakkam, Municipal	Kalyana	Mandapam,	Chennai		
4	Kathivakkam, Municipal	Kalyana	Mandapam,	Chennai		
...
2874		Central Bus	Stand,	Trichy		
2875		Central Bus	Stand,	Trichy		
2876		Central Bus	Stand,	Trichy		
2877		Central Bus	Stand,	Trichy		
2878		Central Bus	Stand,	Trichy		

	Agency					\
0	Tamilnadu	State	Pollution	Control	Board	
1	Tamilnadu	State	Pollution	Control	Board	
2	Tamilnadu	State	Pollution	Control	Board	
3	Tamilnadu	State	Pollution	Control	Board	
4	Tamilnadu	State	Pollution	Control	Board	
...
2874	Tamilnadu	State	Pollution	Control	Board	
2875	Tamilnadu	State	Pollution	Control	Board	
2876	Tamilnadu	State	Pollution	Control	Board	
2877	Tamilnadu	State	Pollution	Control	Board	
2878	Tamilnadu	State	Pollution	Control	Board	

	Type of Location	S02	N02	RSPMorPM10	PM
2.5					
0	Industrial Area	11.0	17.0	55.0	
0.0					
1	Industrial Area	13.0	17.0	45.0	
0.0					

```

2          Industrial Area  12.0  18.0      50.0
0.0
3          Industrial Area  15.0  16.0      46.0
0.0
4          Industrial Area  13.0  14.0      42.0
0.0
...          ...          ...          ...
..
2874  Residential, Rural and other Areas  15.0  18.0      102.0
0.0
2875  Residential, Rural and other Areas  12.0  14.0      91.0
0.0
2876  Residential, Rural and other Areas  19.0  22.0      100.0
0.0
2877  Residential, Rural and other Areas  15.0  17.0      95.0
0.0
2878  Residential, Rural and other Areas  14.0  16.0      94.0
0.0

[2879 rows x 11 columns]

df.duplicated()

0      False
1      False
2      False
3      False
4      False
...
2874   False
2875   False
2876   False
2877   False
2878   False
Length: 2879, dtype: bool

```

Data Normalization

```

from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
df['Values_standardized'] = scaler.fit_transform(df[['RSPMorPM10']])

scaler = StandardScaler()
df['Values_standardized'] = scaler.fit_transform(df[['RSPMorPM10']])
df

```

	Stn Code	Sampling Date	State	City/Town/Village/Area	\
0	38	01-02-2014	Tamil Nadu	Chennai	
1	38	01-07-2014	Tamil Nadu	Chennai	
2	38	21-01-2014	Tamil Nadu	Chennai	

3	38	23-01-2014	Tamil Nadu	Chennai
4	38	28-01-2014	Tamil Nadu	Chennai
...
2874	773	12-03-2014	Tamil Nadu	Trichy
2875	773	12-10-2014	Tamil Nadu	Trichy
2876	773	17-12-2014	Tamil Nadu	Trichy
2877	773	24-12-2014	Tamil Nadu	Trichy
2878	773	31-12-2014	Tamil Nadu	Trichy
Location of Monitoring Station \				
0	Kathivakkam, Municipal Kalyana Mandapam, Chennai			
1	Kathivakkam, Municipal Kalyana Mandapam, Chennai			
2	Kathivakkam, Municipal Kalyana Mandapam, Chennai			
3	Kathivakkam, Municipal Kalyana Mandapam, Chennai			
4	Kathivakkam, Municipal Kalyana Mandapam, Chennai			
...	...			
2874	Central Bus Stand, Trichy			
2875	Central Bus Stand, Trichy			
2876	Central Bus Stand, Trichy			
2877	Central Bus Stand, Trichy			
2878	Central Bus Stand, Trichy			
Agency \				
0	Tamilnadu State Pollution Control Board			
1	Tamilnadu State Pollution Control Board			
2	Tamilnadu State Pollution Control Board			
3	Tamilnadu State Pollution Control Board			
4	Tamilnadu State Pollution Control Board			
...	...			
2874	Tamilnadu State Pollution Control Board			
2875	Tamilnadu State Pollution Control Board			
2876	Tamilnadu State Pollution Control Board			
2877	Tamilnadu State Pollution Control Board			
2878	Tamilnadu State Pollution Control Board			
Type of Location S02 N02 RSPMorPM10 PM				
2.5 \				
0	Industrial Area	11.0	17.0	55.0
NaN				
1	Industrial Area	13.0	17.0	45.0
NaN				
2	Industrial Area	12.0	18.0	50.0
NaN				
3	Industrial Area	15.0	16.0	46.0
NaN				
4	Industrial Area	13.0	14.0	42.0
NaN				
...
...				
2874	Residential, Rural and other Areas	15.0	18.0	102.0


NaN				
2875	Residential, Rural and other Areas	12.0	14.0	91.0
NaN				
2876	Residential, Rural and other Areas	19.0	22.0	100.0
NaN				
2877	Residential, Rural and other Areas	15.0	17.0	95.0
NaN				
2878	Residential, Rural and other Areas	14.0	16.0	94.0
NaN				

	Values_standardized
0	-0.238950
1	-0.557794
2	-0.398372
3	-0.525910
4	-0.653447
...	...
2874	1.259617
2875	0.908889
2876	1.195848
2877	1.036426
2878	1.004542

[2879 rows x 12 columns]

Data Visualisation

```
# Scatter Plot:
plt.scatter(df['S02'], df['RSPMorPM10'], alpha=0.5, color='purple')
plt.title('RSPM/PM10')
plt.xlabel('S02')
plt.ylabel('N02')
plt.show()
```



Data Analysis using different models

Simple Linear Regression

```
cdf = df[['S02', 'N02', 'RSPMorPM10']]
cdf.head(9)
```

	S02	N02	RSPMorPM10
0	11.0	17.0	55.0
1	13.0	17.0	45.0
2	12.0	18.0	50.0
3	15.0	16.0	46.0
4	13.0	14.0	42.0
5	14.0	18.0	43.0
6	12.0	17.0	51.0
7	13.0	16.0	46.0
8	10.0	19.0	50.0

```
msk = np.random.rand(len(df)) < 0.8
train = cdf[msk]
test = cdf[~msk]
```

```
from sklearn import linear_model
regr = linear_model.LinearRegression()
```

```

train = train.dropna()
x = np.asanyarray(train[['S02', 'N02']])
y = np.asanyarray(train[['RSPMorPM10']])
regr.fit(x, y)
# The coefficients
print ('Coefficients: ', regr.coef_)

Coefficients:  [[2.80336749  0.18971444]]

test = test.dropna()
y_hat= regr.predict(test[['S02', 'N02']])
x = np.asanyarray(test[['S02', 'N02']])
y = np.asanyarray(test[['RSPMorPM10']])
print ("Mean Squared Error (MSE) : %.2f"
% np.mean((y_hat - y) ** 2))
# Explained variance score: 1 is perfect prediction
print ('Variance score: %.2f' % regr.score(x, y))

Mean Squared Error (MSE) : 794.16
Variance score: 0.19

C:\Users\savio\anaconda3\Lib\site-packages\sklearn\base.py:457:
UserWarning: X has feature names, but LinearRegression was fitted
without feature names
  warnings.warn(

```

the results show that due to the low variance score this Model is not fit to predict the data of this type

Random forest Algorithm

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt

# Load your dataset from the CSV file
data = pd.read_csv("cpcb_dly_aq_tamil_nadu-2014.csv") # Adjust the filename as needed

# Handle missing values by filling with the mean
data = data.fillna(data.mean())

# Select the relevant features (S02 and N02) and the target variable (RSPM/PM10)
features = data[["S02", "N02"]]
target = data["RSPM/PM10"]

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(features, target,

```

```

test_size=0.2, random_state=42)

# Create and train the Random Forest model
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = rf_model.predict(X_test)

# Calculate model performance metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse:.2f}")
print(f"R-squared (R2) Score: {r2:.2f}")


# Visualize the feature importances
feature_importances = rf_model.feature_importances_
plt.bar(features.columns, feature_importances)
plt.xlabel("Features")
plt.ylabel("Importance")
plt.title("Feature Importances")
plt.show()

```

C:\Users\savio\AppData\Local\Temp\ipykernel_6284\3394349289.py:11: FutureWarning: The default value of numeric_only in DataFrame.mean is deprecated. In a future version, it will default to False. In addition, specifying 'numeric_only=None' is deprecated. Select only valid columns or specify the value of numeric_only to silence this warning.

```
data = data.fillna(data.mean())
```

Mean Squared Error: 775.49
R-squared (R2) Score: 0.25



Regression with the Gradient Boosting algorithm

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt

# Load your dataset from the CSV file
data = pd.read_csv("cpcb_dly_aq_tamil_nadu-2014.csv") # Adjust the
filename as needed

# Handle missing values by filling with the mean
data = data.fillna(data.mean())

# Select the relevant features (SO2 and NO2) and the target variable
(RSPM/PM10)
features = data[["SO2", "NO2"]]
target = data["RSPM/PM10"]

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(features, target,
test_size=0.2, random_state=42)
```



```

# Create and train the Gradient Boosting model
gb_model = GradientBoostingRegressor(n_estimators=100,
random_state=42)
gb_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = gb_model.predict(X_test)

# Calculate model performance metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse:.2f}")
print(f"R-squared (R2) Score: {r2:.2f}")

# Visualize the feature importances
feature_importances = gb_model.feature_importances_
plt.bar(features.columns, feature_importances)
plt.xlabel("Features")
plt.ylabel("Importance")
plt.title("Feature Importances")
plt.show()

```

```

C:\Users\savio\AppData\Local\Temp\ipykernel_6284\2929334411.py:11:
FutureWarning: The default value of numeric_only in DataFrame.mean is
deprecated. In a future version, it will default to False. In
addition, specifying 'numeric_only=None' is deprecated. Select only
valid columns or specify the value of numeric_only to silence this
warning.

```

```


    data = data.fillna(data.mean())

```

```

Mean Squared Error: 684.02
R-squared (R2) Score: 0.34

```



Gaussian Process Regression (GPR)

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.gaussian_process import GaussianProcessRegressor
from sklearn.gaussian_process.kernels import RBF, ConstantKernel as C
import matplotlib.pyplot as plt

# Load your dataset, replace 'your_data.csv' with the actual file path
data = pd.read_csv('cpcb_dly_aq_tamil_nadu-2014.csv')

# Data Preprocessing
# Select relevant columns
data = data[['SO2', 'NO2', 'RSPM/PM10']]

# Check for missing values and handle them if necessary
data.dropna(inplace=True)

# Split the data into features (X) and target (y)
X = data[['SO2', 'NO2']].values
y = data['RSPM/PM10'].values
```

```

# Standardize the features
scaler = StandardScaler()
X = scaler.fit_transform(X)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Define the Gaussian Process Kernel
kernel = C(1.0, (1e-3, 1e3)) * RBF(1.0, (1e-2, 1e2))

# Create and train the Gaussian Process Regressor
gpr = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=10,
random_state=42)
gpr.fit(X_train, y_train)

# Make predictions
y_pred, sigma = gpr.predict(X_test, return_std=True)

# Evaluate the model (you can use various metrics, e.g., RMSE, R-
squared)
from sklearn.metrics import mean_squared_error, r2_score
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
r2 = r2_score(y_test, y_pred)
print(f'RMSE: {rmse:.2f}')
print(f'R-squared: {r2:.2f}')

# Visualize the results
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, color='blue', alpha=0.6)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)],
linestyle='--', color='red', linewidth=2)
plt.xlabel('Actual RSPM/PM10')
plt.ylabel('Predicted RSPM/PM10')
plt.title('Gaussian Process Regression - Actual vs. Predicted')
plt.show()

C:\Users\savio\anaconda3\Lib\site-packages\sklearn\gaussian_process\
kernels.py:429: ConvergenceWarning: The optimal value found for
dimension 0 of parameter k1__constant_value is close to the specified
upper bound 1000.0. Increasing the bound and calling fit again may
find a better value.
  warnings.warn(
C:\Users\savio\anaconda3\Lib\site-packages\sklearn\gaussian_process\
kernels.py:419: ConvergenceWarning: The optimal value found for
dimension 0 of parameter k2__length_scale is close to the specified
lower bound 0.01. Decreasing the bound and calling fit again may find
a better value.
  warnings.warn(

```

RMSE: 33.17
R-squared: -0.05

